



西安交通大学
XI'AN JIAOTONG UNIVERSITY

Systems Engineering Institute
Ministry of Education Key Lab for Intelligent Networks and Network Security

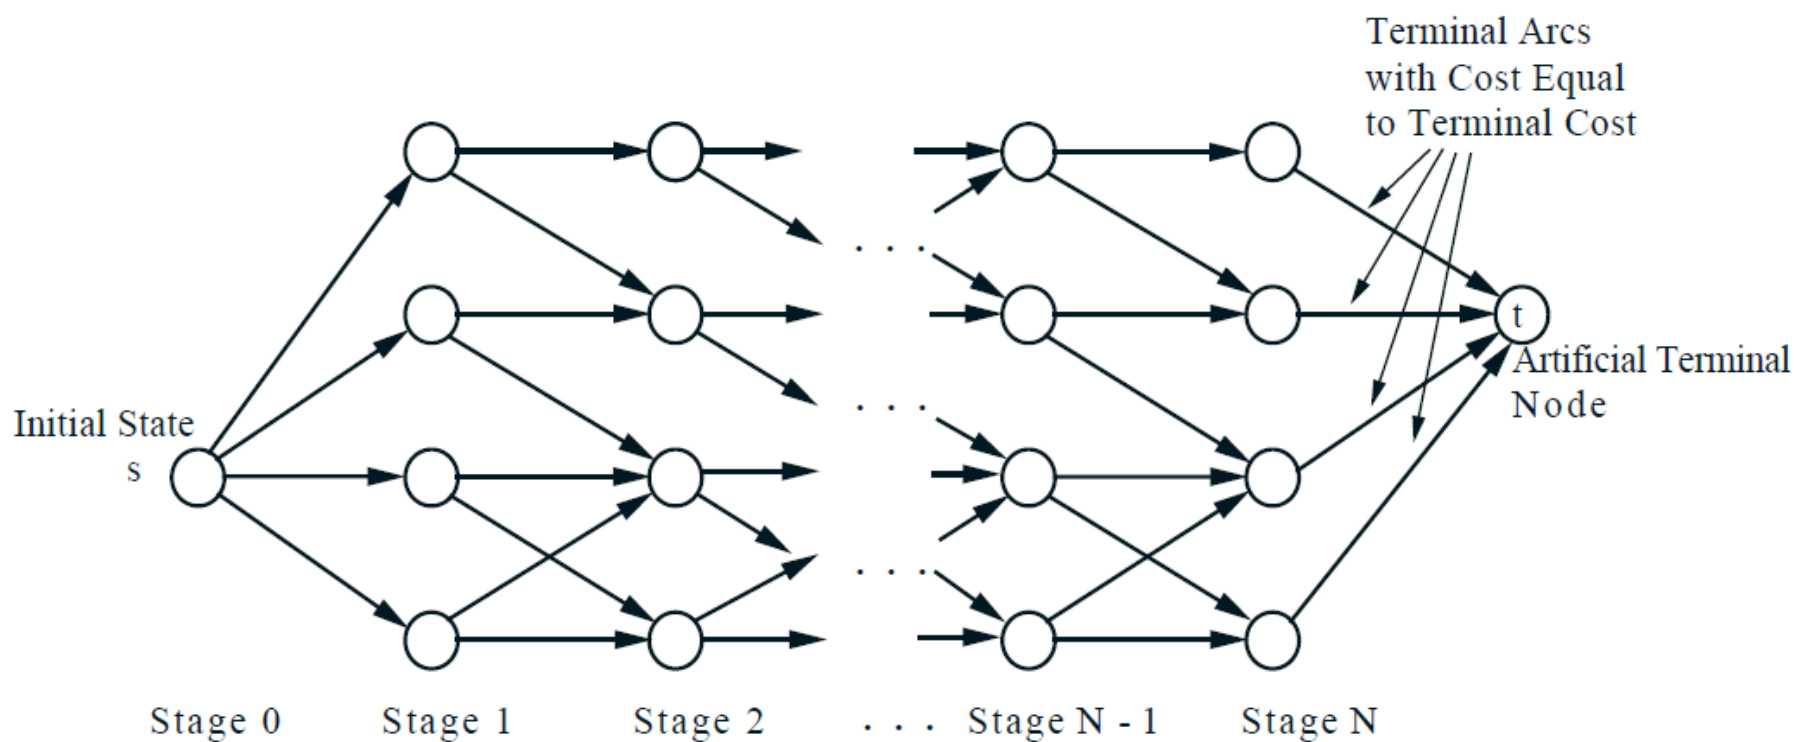
动态规划问题举例 Examples in DP

电信学院·自动化科学与技术系
系统工程研究所
吴江

Outline

- ▶ 确定性定期多阶段决策问题
- ▶ 确定性不定期多阶段决策问题

状态转移图



基本递推方程

$$f_k(x_k) = \min_{u_k} [G(x_k, u_k, k) + f_{k+1}(x_{k+1})]$$

投资分配问题(纯离散问题)

- 某公司计划用40万元投资项目A, B, C. 下表给出了不同投资规模下的预期利润. 试制定最优投资计划

A			B				C			
1	2	3	1	2	3	4	1	2	3	4
20	30	40	10	20	30	40	10	20	30	40
1.8	2.8	3.2	1.2	1.9	2.5	3	0.8	1.6	2.4	3.1

建模

阶段?



投资顺序

状态?



剩余金额

决策?



投资额

转移方程?



$$x_{k+1} = x_k - u_k$$


20	30	40
1.8	2.8	3.2

10	20	30	40
1.2	1.9	2.5	3.0

10	20	30	40
0.8	1.6	2.4	3.1

最优投资方案：

A投30万，B投10万，C投0万元



确定性定期多阶段决策问题

例2: (旅行商问题, Traveling Salesman Problem, TSP)

有 $n + 1$ 个城市, 记为 v_0, v_1, \dots, v_n , 一个推销员从 v_0 出发, 遍访 v_1, \dots, v_n 各恰好一次后再返回 v_0 , 已知从 v_i 到 v_j 的旅费(或路程长度、耗时等)为 $d_{i,j}$, 求最优路线安排。

解: 怎样划分阶段? 按自然时序, 划分为 $n + 1$ 个阶段

怎样定义状态? 状态: 每个阶段/时刻系统所处的状况、态势

状态 (v_i, V) : v_i 为当前时刻所在城市, V 为尚未经过的城市

无后效性? 市集集合(V 中不包含 v_0) 思考: 状态数目? $O(2^n)$

决策 $(v_i, V) \rightarrow (v_j, V \setminus \{v_j\}), v_j \in V$ 决策费用为 $d_{i,j}$

思考: 画状态转移图?

应利用基本方程求解!

确定性定期多阶段决策问题

例2: (旅行商问题, Traveling Salesman Problem, TSP)

状态 (v_i, V) 决策 $(v_i, V) \rightarrow (v_j, V \setminus \{v_j\}), v_j \in V$

怎样列基本方程? 基本方程是关于cost-to-go的递推方程。


$f(v_i, V) = ?$ 从 v_i 出发, 遍访 V 中所有城市各恰好一次, 再回到 v_0 的最短路程长度

状态转移图上求解过程的启示……

边界条件?

$$\begin{cases} f(v_i, \phi) = d_{i,0} & , \quad \forall v_i \neq v_0 \\ f(v_i, V) = \min_{v_j \in V} \{ d_{i,j} + f(v_j, V \setminus \{v_j\}) \} \end{cases}$$

求 $f(v_0, \{v_1, v_2, \dots, v_n\}) = ?$



确定性定期多阶段决策问题

例2: (旅行商问题, Traveling Salesman Problem, TSP)

状态 (v_i, V)

决策 $(v_i, V) \rightarrow (v_j, V \setminus \{v_j\}), v_j \in V$

实例

$$D = \begin{matrix} & \begin{matrix} v_0 & v_1 & v_2 & v_3 \end{matrix} \\ \begin{bmatrix} 0 & 8 & 5 & 6 \\ 6 & 0 & 8 & 5 \\ 7 & 9 & 0 & 5 \\ 9 & 7 & 8 & 0 \end{bmatrix} & \begin{matrix} v_0 \\ v_1 \\ v_2 \\ v_3 \end{matrix} \end{matrix}$$

注意: 非对称TSP


最优解: $v_0 \rightarrow v_2 \rightarrow v_3 \rightarrow v_1 \rightarrow v_0$

P170~171

计算复杂性分析

$$\begin{cases} f(v_i, \phi) = d_{i,0}, & \forall v_i \neq v_0 \\ f(v_i, V) = \min_{v_j \in V} \{d_{i,j} + f(v_j, V \setminus \{v_j\})\} \end{cases}$$

求 $f(v_0, \{v_1, v_2, \dots, v_n\}) = ?$




动态规划 v.s. 非线性(混合整数)规划

- ▶ 确定性定期多阶段决策问题基本上都可以转化为非线性(混合整数)规划问题.
- ▶ 非线性(混合整数)规划问题转化为DP:
 - 最优化原理
 - 无后效性
 - 子问题的重叠性
- ▶ DP求解的原因
 - 全局解v.s.局部解
 - 中间信息
 - 求解效率

基本递推方程

$$f_k(x_k) = \min_{u_k} [G(x_k, u_k, k) + f_{k+1}(x_{k+1})]$$



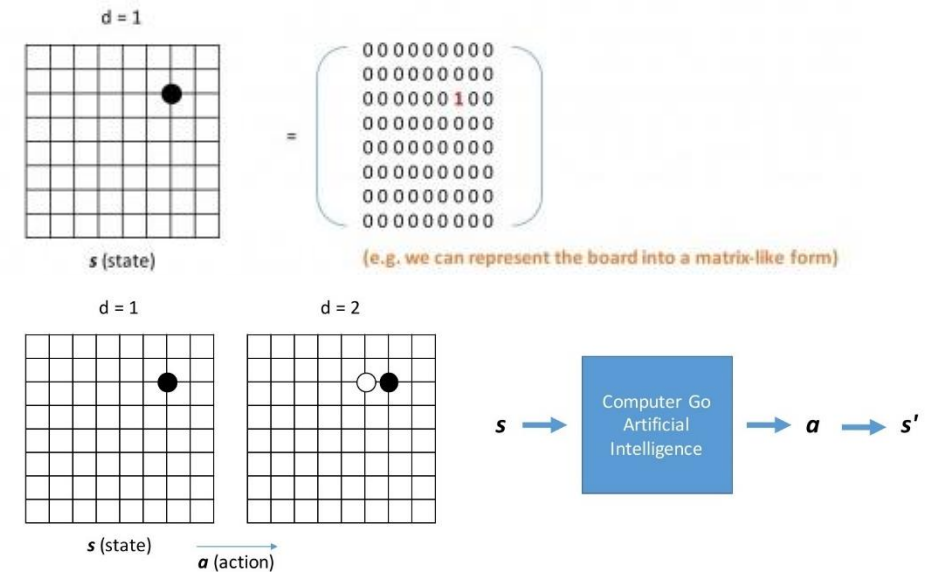
$$\pi^*(s) = \arg_a \min [r(s, a) + V^*(\delta(s, a))]$$

Mastering the Game of Go with Deep Neural Networks and Tree Search

(Nature 529, 484–489, 28 January 2016)



Computer Go AI – Definition



$$\pi^*(s) = \arg_a \max [r(s, a) + V^*(\delta(s, a))]$$

$$|S| = 3^{361} \quad |A_k| = 361 - 2(k-1) \quad |\Omega| = 361 * 359 * 357 * \dots$$

Reducing "act candidates"

Current Board

```

00 000 0000
00 000 1000
0-100 1-1100
0 1 00 1-1000
00 00-10000
00 000 0000
0-1000 0000
00 000 0000
    
```

Prediction Model

Next Action

```

0000000000
0000000000
0000000000
0000000000
0000000000
0000001000
0000000000
0000000000
0000000000
    
```

s

$f: s \rightarrow a$

a

Expert Moves Imitator Model
(w/ CNN)

30,000,000 < s, a >

Updated Model
ver 1.3

vs

Updated Model
ver 1.7

30,000,000 < s, a >

Current Board

```

00 000 0000
00 000 1000
0-100 1-1100
0 1 00 1-1000
00 00-10000
00 000 0000
0-1000 0000
00 000 0000
    
```

Deep Learning
(13 Layer CNN)

```

000000 000
000000 000
000000 000
000000 000
000000.20.100
000000.40.200
000000.1 000
000000 000
000000 000
000000 000
    
```

Next Action

```

0000000000
0000000000
0000000000
0000000000
0000000000
0000001000
0000000000
0000000000
0000000000
    
```

s

$g: s \rightarrow p(a|s)$

$p(a|s)$

argmax

a

Reducing "act candidates"

ICLR 2015


Feature	Planes	Description
Black / white / empty	3	Stone colour
Liberties	4	Number of liberties (empty adjacent points)
Liberties after move	6	Number of liberties after this move is played
Legality	1	Whether point is legal for current player
Turns since	5	How many turns since a move was played
Capture size	7	How many opponent stones would be captured
Ladder move	1	Whether a move at this point is a successful ladder capture
KGS rank	9	Rank of current player

Extended Data Table 2 | Input features for neural networks

Nature 2016

Feature	# of planes	Description
Stone colour	3	Player stone / opponent stone / empty
Ones	1	A constant plane filled with 1
Turns since	8	How many turns since a move was played
Liberties	8	Number of liberties (empty adjacent points)
Capture size	8	How many opponent stones would be captured
Self-atari size	8	How many of own stones would be captured
Liberties after move	8	Number of liberties after this move is played
Ladder capture	1	Whether a move at this point is a successful ladder capture
Ladder escape	1	Whether a move at this point is a successful ladder escape
Sensibleness	1	Whether a move is legal and does not fill its own eyes
Zeros	1	A constant plane filled with 0
Player color	1	Whether current player is black

Feature planes used by the policy network (all but last feature) and value network (all features).



Current Board

```

00 000 0000
00 000 1000
0-100 1-1100
0 1 00 1-1000
00 00-10000
00 000 0000
0-1000 0000
00 000 0000

```

Deep Learning
(13 Layer CNN)

Next Action

```

000000 000
000000 000
000000 000
000000.20.100
000000.40.200
000000.1 000
000000 000
000000 000

```

s


$g: s \rightarrow p(a|s)$

$p(a|s)$

argmax




a

Board Evaluation

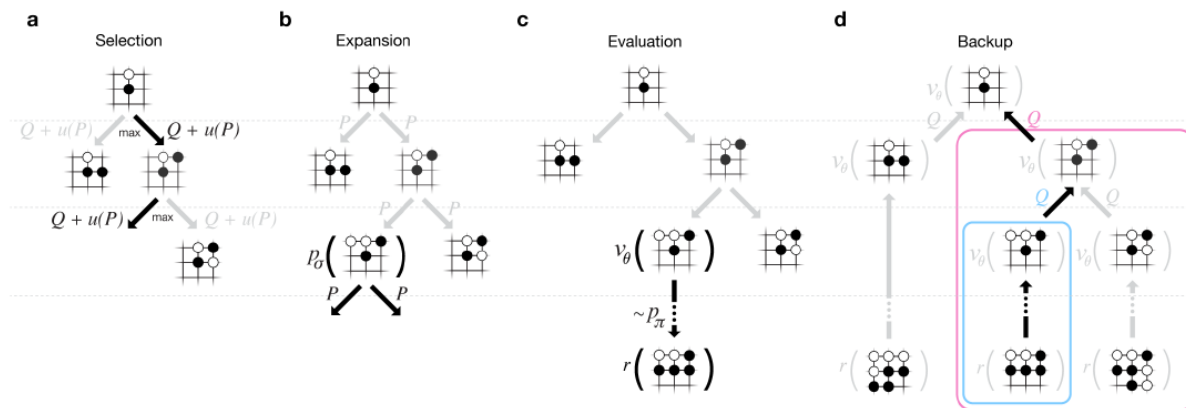


Instead of simulating until


Cost to go?




Monte-Carlo tree search



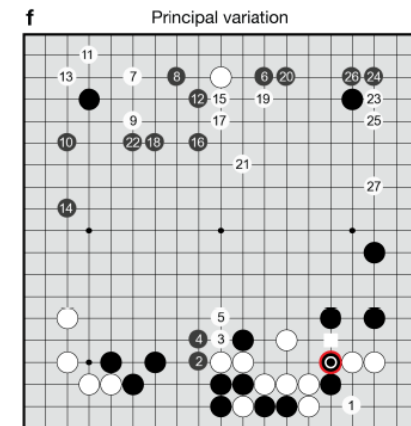





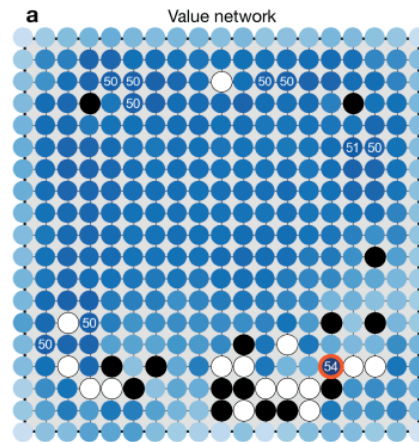
How AlphaGo selected its move



Bread reduction




Depth reduction



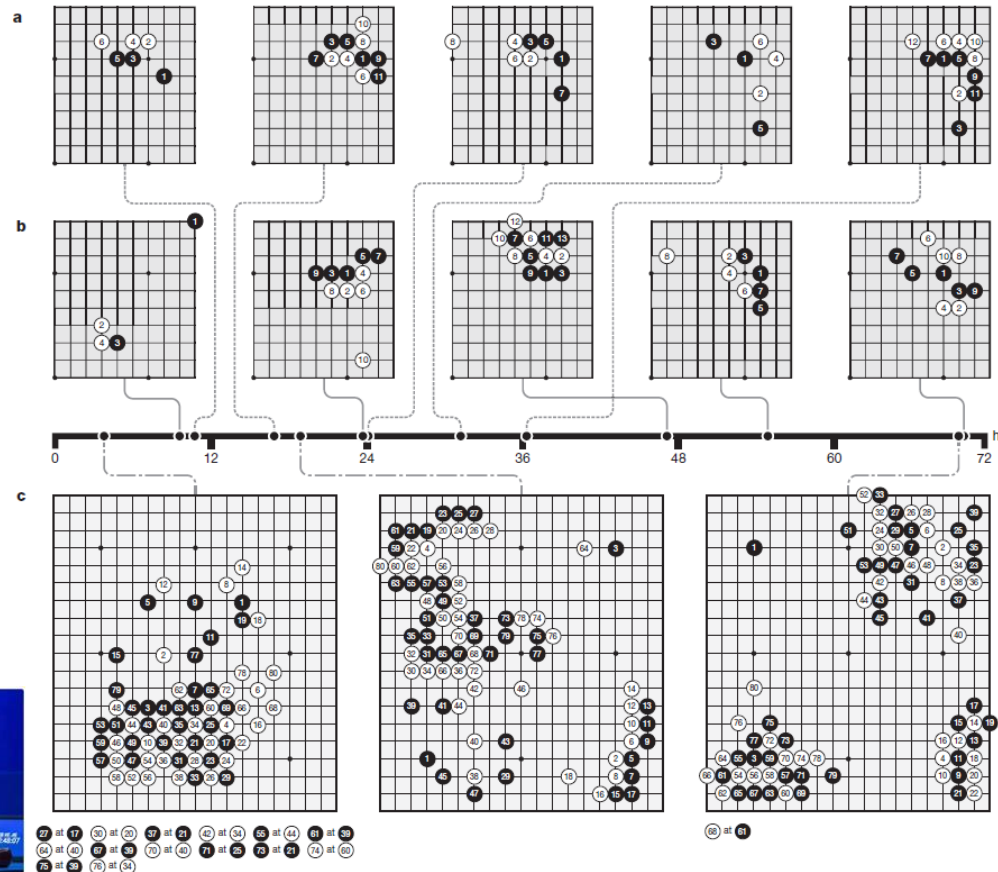

Nature 2017: Mastering the game of Go without human knowledge

1. without any human data
2. only stones as input features
3. single neural network
4. without any Monte Carlo rollouts

a Self-play



b Neural network training



Challenges of Real-World Reinforcement Learning

---- ICML 19'

1	Training off-line from the fixed logs of an external behavior policy.
2	Learning on the real system from limited samples .
3	High -dimensional continuous state and action spaces.
4	Safety constraints that should never or at least rarely be violated.
5	Tasks that may be partially observable , alternatively viewed as non-stationary or stochastic.
6	Reward functions that are unspecified, multi-objective, or risk-sensitive.
7	System operators who desire explainable policies and actions.
8	Inference that must happen in real-time at the control frequency of the system.
9	Large and/or unknown delays in the system actuators, sensors, or rewards.